


Link github: https://github.com/NguyenHuynhThaoNhu/MKTG5883.N22_19521970

I. Classification

```
[ ] from sklearn import datasets
import matplotlib.pyplot as plt
digits = datasets.load_digits()
```

```
[ ] plt.imshow(digits.images[1010], cmap = plt.cm.gray_r, interpolation='nearest')
```



```
[ ] from sklearn.model_selection import train_test_split
X = digits.data
y = digits.target


X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2,
                                                    random_state = 42, stratify = y)
```

```
from sklearn.neighbors import KNeighborsClassifier
import numpy as np
knn = KNeighborsClassifier(n_neighbors = 3)
knn.fit(X_train, y_train)
print("Accuracy: {}".format(knn.score(X_test, y_test)))
```

Accuracy: 0.9861111111111112

```
for i, k in enumerate(neighbors):
    knn = KNeighborsClassifier(n_neighbors = k)
    knn.fit(X_train, y_train)
    train_accuracy[i] = knn.score(X_train, y_train)
    test_accuracy[i] = knn.score(X_test, y_test)

plt.title('k-NN: varying number of neighbors')
plt.plot(neighbors, test_accuracy, label = 'Testing accuracy')
plt.plot(neighbors, train_accuracy, label = 'Training accuracy')
plt.legend()
plt.xlabel('number of neighbors')
plt.ylabel('accuracy')
plt.show()
```



```

[8] from __future__ import print_function
import torch
import torch.nn as nn
import torch.nn.functional as F
from torch.autograd import Variable

[9] from torchvision import datasets, transforms
mnist = datasets.MNIST(root='.', train=True, download=True)

Downloading http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz
Downloading http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz to ./MNIST/raw/train-images-idx3-ubyte.gz
100% [#####] 9912422/9912422 [00:00<00:00, 80603544.45it/s]
Extracting ./MNIST/raw/train-images-idx3-ubyte.gz to ./MNIST/raw

Downloading http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz
Downloading http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz to ./MNIST/raw/train-labels-idx1-ubyte.gz
100% [#####] 20891/20891 [00:00<00:00, 61857055.07it/s]Extracting ./MNIST/raw/train-labels-idx1-ubyte.gz to ./MNIST/raw

Downloading http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz
Downloading http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz to ./MNIST/raw/t10k-images-idx3-ubyte.gz
100% [#####] 1648877/1648877 [00:00<00:00, 26040414.47it/s]
Extracting ./MNIST/raw/t10k-images-idx3-ubyte.gz to ./MNIST/raw

Downloading http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz
Downloading http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz to ./MNIST/raw/t10k-labels-idx1-ubyte.gz
100% [#####] 4542/4542 [00:00<00:00, 9597243.71it/s]
Extracting ./MNIST/raw/t10k-labels-idx1-ubyte.gz to ./MNIST/raw

```

```

[10] print("Number of training examples", mnist.train_data.shape)
print("Image information", mnist[0])


Number of training examples torch.Size([60000, 28, 28])
Image information (PIL.Image.Image mode='L' size=28x28 at 0x7f051e53c220, 5)
/usr/local/lib/python3.9/dist-packages/torchvision/datasets/mnist.py:75: UserWarning: train_data has been renamed data
warnings.warn("train_data has been renamed data")

```

```

import matplotlib.pyplot as plt
#matplotlib inline
plt.imshow(mnist[0][0])

```



```

[12] class Net(nn.Module):
def __init__(self):
super(Net, self).__init__()

self.fully = nn.Sequential(
nn.Linear(28*28, 10)
)
def forward(self, x):
x = x.view([-1, 28*28])
x = self.fully(x)
x = F.log_softmax(x, dim = 1)
return x

```

```

[22] from scipy.spatial import transform
train_loader = torch.utils.data.DataLoader(datasets.MNIST(root='.', train = True, transform = transforms.Compose([transforms.ToTensor()])),batch_size = 64, shuffle = True)
test_loader = torch.utils.data.DataLoader(datasets.MNIST(root='.', train = False, transform = transforms.Compose([transforms.ToTensor()])),batch_size = 1, shuffle = True)

```

```

[41] def train():
learning_rate = 1e-3
num_epochs = 3

net = Net()
optimizer = torch.optim.Adam(net.parameters(), lr=learning_rate)

for epoch in range(num_epochs):
for batch_idx, (data, target) in enumerate(train_loader):
output = net(data)

loss = F.nll_loss(output, target)
optimizer.zero_grad()
loss.backward()
optimizer.step()

if batch_idx % 100 == 0:
print('Epoch = %f. Batch = %s. Loss = %s' % (epoch, batch_idx, loss.item()))

return net

```

```
net = train()

Epoch = 0.000000. Batch = 0. Loss = 2.3549129962921143
Epoch = 0.000000. Batch = 100. Loss = 0.7764542102813721
Epoch = 0.000000. Batch = 200. Loss = 0.5944685881787189
Epoch = 0.000000. Batch = 300. Loss = 0.3326780972454071
Epoch = 0.000000. Batch = 400. Loss = 0.4065855145454407
Epoch = 0.000000. Batch = 500. Loss = 0.37920111083984375
Epoch = 0.000000. Batch = 600. Loss = 0.3410077287387848
Epoch = 0.000000. Batch = 700. Loss = 0.5293893814086914
Epoch = 0.000000. Batch = 800. Loss = 0.418151319026947
Epoch = 0.000000. Batch = 900. Loss = 0.386769046226293
Epoch = 1.000000. Batch = 0. Loss = 0.2576746344566345
Epoch = 1.000000. Batch = 100. Loss = 0.2363649159669876
Epoch = 1.000000. Batch = 200. Loss = 0.38338365896224976
Epoch = 1.000000. Batch = 300. Loss = 0.3391488492468861
Epoch = 1.000000. Batch = 400. Loss = 0.25349998474121894
Epoch = 1.000000. Batch = 500. Loss = 0.16868181068390472
Epoch = 1.000000. Batch = 600. Loss = 0.4312569499615888
Epoch = 1.000000. Batch = 700. Loss = 0.16464556753635486
Epoch = 1.000000. Batch = 800. Loss = 0.29552900791168213
Epoch = 1.000000. Batch = 900. Loss = 0.31031396985854016
Epoch = 2.000000. Batch = 0. Loss = 0.24668877361216125
Epoch = 2.000000. Batch = 100. Loss = 0.46185773611060726
Epoch = 2.000000. Batch = 200. Loss = 0.24814119935035706
Epoch = 2.000000. Batch = 300. Loss = 0.18867127597332
Epoch = 2.000000. Batch = 400. Loss = 0.2038491666316986
Epoch = 2.000000. Batch = 500. Loss = 0.21730600178756714
Epoch = 2.000000. Batch = 600. Loss = 0.3387146430015564
Epoch = 2.000000. Batch = 700. Loss = 0.48744587779845185
Epoch = 2.000000. Batch = 800. Loss = 0.387571172142334
Epoch = 2.000000. Batch = 900. Loss = 0.2591584752635956

[44]: net.eval()
test_loss = 0
correct = 0
total = 0

for data, target in test_loader:
    total += 1
    output = net(data)
    pred = output.max(1, keepdim = True)[1]
    correct += target.eq(pred.view_as(target)).sum()

print("Correct out of %s" % total, correct.item())
print("Percentage accuracy", correct.item()*100/10000.)

Correct out of 10000 9228
Percentage accuracy 92.28
```

II. Linear Regression

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt

df = pd.read_csv('gapminder.csv')
ax = sns.heatmap(df.corr(), square = True, cmap = 'RDYlgn')
plt.show()

thon-input-47-9dd44cb24381>: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric_only to silence
sns.heatmap(df.corr(), square = True, cmap = 'RDYlgn')
```

```
[48]: from sklearn.linear_model import LinearRegression, LogisticRegression
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error
from sklearn.model_selection import train_test_split

[51]: x_fertility = df['fertility'].values.reshape(-1,1)
y_life = df['life'].values.reshape(-1,1)
prediction_space = np.linspace(min(x_fertility),max(x_fertility)).reshape(-1,1)
x_train, x_test, y_train, y_test = train_test_split(x_fertility, y_life, test_size=0.3, random_state=42)

[52]: reg = LinearRegression()


reg.fit(x_train, y_train)
y_predict = reg.predict(prediction_space)
print(reg.score(x_fertility, y_life))

0.6162438752151917
```

```

plt.scatter(x_fertility, y_life, color = 'blue')
plt.plot(prediction_space, y_predict, color='black', linewidth = 3)
plt.ylabel('Life expectancy')
plt.xlabel('fertility')
plt.show()

```



```

[61] features = pd.read_csv('gapminder.csv')
df = pd.read_csv('gapminder.csv')
del features['life']
del features['Region']

y_life = df['life'].values.reshape(-1,1)

#Create training and test sets
x_train, x_test, y_train, y_test = train_test_split(features, y_life, test_size = 0.3, random_state=42)

#Create the regression model: reg_all
reg_all = LinearRegression()

#Fit the regression to the training data
reg_all.fit(x_train, y_train)

#Print accuracy
print(reg_all.score(features, y_life))

0.8914651485793176

```

Linear Regression using Pytorch

```


[62] import matplotlib.pyplot as plt
%matplotlib inline
import numpy as np

```

```

# N = 10 number of data points
N = 10
c = 1
x = np.linspace(0, 2 * np.pi, N)
y = 0.5 * c + np.random.normal(0, 1, x.shape)
plt.figure()
plt.plot(x, y, 'x')
plt.xlabel('x')
plt.ylabel('y')
plt.title('2D data (data = 10)')
plt.show()

```



```

[77] import torch

```

Dataset

```

[80] from torch.utils.data import Dataset

class MyDataset(Dataset):
    def __init__(self, x, y):
        self.x = x
        self.y = y

    def __len__(self):
        return len(self.x)

    def __getitem__(self, idx):
        sample = {
            'feature': torch.tensor([1, self.x[idx]]),
            'label': torch.tensor([self.y[idx]])
        }
        return sample

```

```

[81] dataset = MyDataset(x, y)

for i in range(len(dataset)):
    sample = dataset[i]
    print(i, sample['feature'], sample['label'])

0 tensor([1., 0.], dtype=torch.float64) tensor([1.4156], dtype=torch.float64)
1 tensor([1.0000, 0.6981], dtype=torch.float64) tensor([1.8023], dtype=torch.float64)
2 tensor([1.0000, 1.3963], dtype=torch.float64) tensor([1.9613], dtype=torch.float64)
3 tensor([1.0000, 2.0944], dtype=torch.float64) tensor([3.1067], dtype=torch.float64)
4 tensor([1.0000, 2.7925], dtype=torch.float64) tensor([3.6164], dtype=torch.float64)
5 tensor([1.0000, 3.4907], dtype=torch.float64) tensor([4.9668], dtype=torch.float64)
6 tensor([1.0000, 4.1888], dtype=torch.float64) tensor([4.8568], dtype=torch.float64)
7 tensor([1.0000, 4.8869], dtype=torch.float64) tensor([5.7871], dtype=torch.float64)
8 tensor([1.0000, 5.5851], dtype=torch.float64) tensor([6.0412], dtype=torch.float64)
9 tensor([1.0000, 6.2832], dtype=torch.float64) tensor([6.4988], dtype=torch.float64)

```

▼ Dataloader

```
[99] from torch.utils.data import DataLoader

dataset = MyDataset(x, y)
batch_size = 4
shuffle = True
num_workers = 4
dataloader = DataLoader(dataset, batch_size=batch_size, shuffle=shuffle, num_workers=num_workers)

/usr/local/lib/python3.9/dist-packages/torch/utils/data/dataloader.py:561: UserWarning: This DataLoader will create 4 worker processes in total. Our suggested max number of worker in current system is 2, which is smaller than what this
warnings.warn(_create_warning_msg(

import pprint as pp
for i_batch, samples in enumerate(dataloader):
    print('\nbatch# = %s' % i_batch)
    print('sample:')
    pp.pprint(samples)

batch# = 0
sample:
('feature': tensor([[1.0000, 0.6981],
                    [1.0000, 0.0000],
                    [1.0000, 4.8889],
                    [1.0000, 2.7925]], dtype=torch.float64),
 'label': tensor([[1.8823],
                  [1.4156],
                  [3.7871],
                  [3.6544]], dtype=torch.float64))

batch# = 1
sample:
('feature': tensor([[1.0000, 2.8944],
                    [1.0000, 6.2832],
                    [1.0000, 5.5851],
                    [1.0000, 4.1888]], dtype=torch.float64),
 'label': tensor([[3.1867],
                  [6.4988],
                  [6.0422],
                  [4.8560]], dtype=torch.float64))

batch# = 2
sample:
('feature': tensor([[1.0000, 3.4907],
                    [1.0000, 1.3963]], dtype=torch.float64),
 'label': tensor([[4.9668],
                  [1.9613]], dtype=torch.float64))
```

▼ Model

```
[101] import torch.nn as nn
import torch.nn.functional as F
class MyModel(nn.Module):
    def __init__(self, input_dim, output_dim):
        super(MyModel, self).__init__()
        self.linear = nn.Linear(input_dim, output_dim)

    def forward(self, x):
        out = self.linear(x)
        return out

[102] input_dim = 2
output_dim = 1
model = MyModel(input_dim, output_dim)

[103] cost = nn.MSELoss
```

▼ Minimizing the cost function

```
num_epochs = 10
lr_rate = 0.01
optimizer = torch.optim.SGD(model.parameters(), lr=lr_rate)

dataset = MyDataset(x, y)
batch_size = 4
training_sample_generator = DataLoader(dataset, batch_size=batch_size, shuffle=shuffle, num_workers=num_workers)

for epoch in range(num_epochs):
    print('Epoch = %s' % epoch)
    for batch_i, samples in enumerate(training_sample_generator):
        predictions = model(samples['feature'])
        error = cost(predictions, samples['label'])
        print('\tbatch = %s, Error = %s' % (batch_i, error.item()))
        optimizer.zero_grad()
        error.backward()
        optimizer.step()
```

↳ Lets see how well the model has learnt the data

```
1 x_for_plotting = np.linspace(0, 2*np.pi, 1000)
2 design_matrix = torch.tensor(np.vstack([np.ones(x_for_plotting.shape), x_for_plotting])).T, dtype = torch.float32)
3 print('Design matrix shape: ', design_matrix.shape)
4
5 y_for_plotting = model.forward(design_matrix)
6 print('y_for_plotting shape: ', y_for_plotting.shape)
7
8 Design matrix shape: torch.Size([1000, 2])
9 y_for_plotting shape: torch.Size([1000, 1])
```

```
1 plt.figure()
2 plt.plot(x, y, 'o')
3 plt.plot(x_for_plotting, y_for_plotting.data.numpy(), 'r-')
4 plt.xlabel('x')
5 plt.ylabel('y')
6 plt.title('2D data(#data = %d)' % N)
7 plt.show()
```

